

# Computação em Larga Escala

## CUDA Programming

Eurico Pedrosa

António Rui Borges

Universidade de Aveiro - DETI

2025-05-01

# CUDA

---

- CUDA is a **general-purpose parallel computing platform** and **programming model** developed by **NVIDIA**.
- It leverages the **parallel processing power** of NVIDIA **GPUs** to efficiently solve **computation-intensive problems**.

- The CUDA platform is accessible via:
  - ▶ **CUDA-accelerated libraries**
  - ▶ **Compiler directives**
  - ▶ **Application Programming Interfaces (APIs)**
  - ▶ **Extensions** to industry-standard languages such as **C, C++, Fortran, Java, and Python**
- **CUDA C** is an **extension of ANSI C** that includes:
  - ▶ Language features for **heterogeneous programming**
  - ▶ APIs to manage **devices, memory, and other GPU tasks**

GPU Computing Applications						
Libraries and Middleware						
cuDNN TensorRT	cuFFT cuBLAS cuRAND cuSPARSE	CULA MAGMA	Thrust NPP	VSIPL SVM OpenCurrent	PhysX OptiX iRay	MATLAB Mathematica
Programming Languages						
C	C++	Fortran	Java Python Wrappers	DirectCompute	Directives (e.g. OpenACC)	
CUDA-Enabled NVIDIA GPUs						
NVIDIA Ampere Architecture (compute capabilities 8.x)					Tesla A Series	
NVIDIA Turing Architecture (compute capabilities 7.x)			GeForce 2000 Series	Quadro RTX Series	Tesla T Series	
NVIDIA Volta Architecture (compute capabilities 7.x)	DRIVE/JETSON AGX Xavier			Quadro GV Series	Tesla V Series	
NVIDIA Pascal Architecture (compute capabilities 6.x)	Tegra X2		GeForce 1000 Series	Quadro P Series	Tesla P Series	
						

Figure 1: GPU Computing Applications. Source: CUDA C Programming Guide, NVIDIA

- CUDA provides **two API levels** for managing the GPU and organizing threads:
  - **CUDA Driver API** ([docs.nvidia.com/cuda/cuda-driver-api](https://docs.nvidia.com/cuda/cuda-driver-api))
  - **CUDA Runtime API** ([docs.nvidia.com/cuda/cuda-runtime-api](https://docs.nvidia.com/cuda/cuda-runtime-api))

- **Driver API:**
  - A **low-level API** that gives **fine-grained control** over GPU behavior
  - **Harder to program**, but offers **more flexibility**
- **Runtime API:**
  - A **high-level API** built **on top of the Driver API**
  - Easier to use; each runtime function maps to one or more driver-level operations

- **Key notes:**
  - ▶ There is **no significant performance difference** between the two APIs
  - ▶ **Thread organization** and **memory usage** impact performance more than API choice
  - ▶ **Mutually exclusive**: you must use **either** the Driver API **or** the Runtime API in your code – **no mixing allowed**

# Compilation Model

- The **CUDA nvcc compiler** separates **host code** and **device code** during compilation
- **Host code:**
  - Written in **standard C/C++**
  - Compiled using a **regular C/C++ compiler**
- **Device code:**
  - Written in **CUDA C** with special **keywords** to label **kernels** (data-parallel functions)
  - Compiled by **nvcc** for execution on the **GPU**

# Compilation Model


- During the **linking stage**:
  - **CUDA runtime** or **driver libraries** are linked in
  - Enables **kernel calls** and **explicit GPU management**
- The **CUDA Toolkit** includes:
  - The **nvcc compiler**
  - **Mathematical libraries**
  - **Debugging and performance optimization tools**

- CUDA exposes three core abstractions through minimal language extensions:
  - **Hierarchy of thread groups**
  - **Shared memory**
  - **Barrier synchronization**
- These abstractions support:
  - **Fine-grained data parallelism** and **thread parallelism**
  - Nested within **coarse-grained data parallelism** and **task parallelism**

- **Problem decomposition** strategy:
  - A problem is partitioned into **subproblems**, each solved in parallel by **thread blocks**
  - Each subproblem is further subdivided and solved cooperatively by **threads within a block**

- Benefits:
  - ▶ Preserves **language expressiveness** through intra-block **thread cooperation**
  - ▶ Enables **scalable execution**:
    - Thread blocks are **independent units of execution**
    - The **runtime system** schedules blocks on **any available multiprocessor**, in **any order**
    - Allows execution on **any GPU**, regardless of the number of multiprocessors

- A **GPU** consists of an array of **Streaming Multiprocessors (SMs)**
- Internally organized as a **MIMD** (Multiple Instruction, Multiple Data) topology of **SIMD** (Single Instruction, Multiple Data) processors
- A **multithreaded CUDA program** is divided into **blocks of threads**
  - **Thread blocks** execute **independently** of each other
  - This allows for **scalable execution** across different GPU hardware
- **Scalability benefit:**
  - A GPU with **more multiprocessors** can execute **more blocks concurrently**
  - This leads to **faster execution** without requiring changes to the code



## Automatic Scalability

Source: CUDA C Programming Guide, NVIDIA

## Typical CUDA Program Structure

1. **Allocate GPU memory**
2. **Copy data from CPU (host) memory to GPU (device) memory**
3. **Invoke CUDA kernel** to perform program-specific computation
4. **Copy results back from GPU memory to CPU memory**
5. **Free (destroy) GPU memory**

# Program Structure

CUDA

```
#include "../common/common.h"
#include <stdio.h>
/* A simple introduction to programming in CUDA. This program
 * prints "Hello World from GPU!" from 10 CUDA threads running
 * on the GPU.
*/
__global__ void helloFromGPU()
{
    printf("Hello World from GPU!\n");
}

int main(int argc, char **argv)
{
    printf("Hello World from CPU!\n");
    helloFromGPU<<<1, 10>>>();
    CHECK(cudaDeviceReset());
    return 0;
}
```

# CUDA Programming Model

---


- **CUDA C** extends C by allowing the definition of **kernels** – special functions executed **concurrently by N CUDA threads**
  - In contrast to regular C functions, which execute only once
- A **kernel** is defined using the **`_global_`** specifier
- A kernel is launched with an **execution configuration**:
  - Syntax: `<<<gridDim, blockDim>>>`
  - Defines a **grid** of **thread blocks** running concurrently

- Each thread is assigned a **unique thread ID** within its block:
  - ▶ Accessible via the built-in variable **threadIdx**
- Each thread has:
  - ▶ Its own **registers** and **private memory**
  - ▶ Access to **input data**
  - ▶ Responsibility to **compute and output results**

- **threadIdx** is a **3-component vector**:
  - ▶ Threads can be indexed in **1D, 2D, or 3D**
  - ▶ Forms a **1D, 2D, or 3D thread block**
  - ▶ Ideal for computations over **vectors, matrices, or volumes**
- **Thread ID computation** (assuming block dimensions Dx, Dy, Dz):
  - ▶ **1D block**:
    - Thread ID = x
  - ▶ **2D block**:
    - Thread ID = x + y × Dx
  - ▶ **3D block**:
    - Thread ID = x + y × Dx + z × Dx × Dy

- **Hardware constraints:**
  - A **thread block** can have up to **1024 threads** (current GPUs)
  - All threads in a block **reside on the same SM (Streaming Multiprocessor)**
  - Threads **share memory resources** (registers, shared memory) of that SM
- **Scalability:**
  - A **kernel** can be launched with **multiple blocks**
  - Total number of threads = **threads per block × number of blocks**

- Thread **blocks** are organized into a **1D, 2D, or 3D grid**
- The number of **blocks per grid** is usually determined by:
  - ▶ The **size of the data set**, or
  - ▶ The **number of available GPU processors** (can exceed both)
- Execution configuration (`<<<...>>>`) supports:
  - ▶ **Thread block dimensions** and **grid dimensions** of type `int` or `dim3`
- Within a kernel:
  - ▶ Each block is identified by **blockIdx**
    - A **1D, 2D, or 3D index** depending on grid shape
  - ▶ The size of each block is given by **blockDim**



Source: CUDA C Programming Guide, NVIDIA

# Thread Block Independence & Synchronization

CUDA Programming Model

- Thread blocks must execute independently:
  - Can be scheduled **in any order, in parallel, or in series**
  - Enables **scalability** across GPUs with varying numbers of cores
- This independence allows blocks to be **assigned dynamically to any core**, supporting scalable parallel execution


# Thread Block Independence & Synchronization

CUDA Programming Model

- Threads within a **block** can **cooperate** through:
  - ▶ **Shared memory** (fast, low-latency memory near the processor core)
  - ▶ **Barrier synchronization** using `__syncthreads()`
- `__syncthreads()`:
  - ▶ Acts as a **barrier** – all threads in the block must reach it before any can proceed
  - ▶ Enables **coordinated memory access**
  - ▶ Expected to be a **lightweight operation**, like L1 cache synchronization in CPUs

- CUDA threads access multiple memory spaces during execution:
  - Local memory: private to each thread
  - Shared memory: accessible by all threads in a block, with block lifetime
  - Global memory: accessible by all threads across all blocks
  - Constant memory: read-only, accessible by all threads
  - Texture memory: read-only, supports special addressing and data filtering
- Global, constant, and texture memory:
  - Optimized for different usage patterns
  - Are persistent across kernel launches within the same application

- **Host vs. Device memory:**
  - ▶ CUDA assumes **separate DRAM** for **host (CPU)** and **device (GPU)**
  - ▶ Requires explicit **memory management** via the **CUDA runtime API**, including:
    - **Device memory allocation/deallocation**
    - **Data transfers** between **host memory** and **device memory**




## Memory Hierarchy

Source: CUDA C Programming Guide, NVIDIA

### Host–Device Memory and Unified Memory in CUDA

- CUDA assumes **separate memory spaces** for:
  - ▶ **Host (CPU) → Host memory**
  - ▶ **Device (GPU) → Device memory**
- Programs must manage:
  - ▶ **Global, constant, and texture memory spaces** via the **CUDA Runtime API**
  - ▶ This includes:
    - **Device memory allocation/deallocation**
    - **Explicit data transfers** between **host** and **device**

- **Unified Memory** (also called **Managed Memory**):
  - ▶ Provides a **single, coherent memory space** shared between **CPU and GPU**
  - ▶ Enables:
    - **Automatic data migration**
    - **Oversubscription** of device memory
    - **Simplified programming** – no need for manual memory copying
    - Easier **application porting** without explicit host–device mirroring



## Heterogeneous Programming

Source: CUDA C Programming Guide, NVIDIA

Qualifier	Executed On	Callable From	Notes
<code>__global__</code>	Device	Host	Must return <code>void</code> ; launches a kernel
<code>__device__</code>	Device	Device	Used for device-only functions
<code>__host__</code>	Host	Host	Optional; default for host code

- `__device__` and `__host__` can be **used together**:
  - ▶ This compiles the function for **both host and device use**
- Common usage:
  - ▶ **Kernel functions** use `__global__`
  - ▶ **Helper functions** shared by host and device may use `__host__ __device__`

CUDA **kernels** are special functions with several important restrictions:

- Must **only access device memory**
- Must have a **void return type**
- **Variable argument lists** (e.g., ...) are **not supported**
- **Static variables** are **not allowed**
- **Function pointers** are **not supported**
- Kernel launches are **asynchronous**
  - ▶ Control returns to the host **before** the kernel finishes execution

## Suggested Reading

---

- CUDA C Programming Guide, NVIDIA
  - ▶ (<https://docs.nvidia.com/cuda/cuda-c-programming-guide/>)